Portrait Optimization Via Deep Learning Networks:

A Framework for Context-Aware Photography Recommendations

Shriya Nair & Bucky Hayes

Institute for Computing in Research Summer 2025


Photo-aid

Photo-aid

Attribute: blur Epochs: 100 Batch-size: 100 LR: 0.00278
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 62.35502004623413 seconds
Maximum Accuracy: 0.9735000133514404

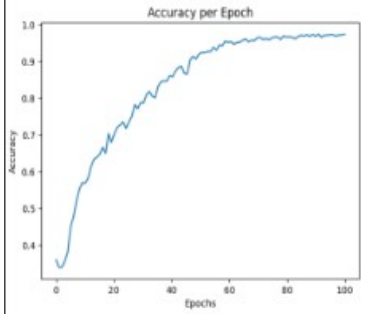

Attribute: occlusion Epochs: 100 Batch-size: 50 LR: 0.0032585
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 90.07953882217407 seconds
Maximum Accuracy: 0.9549999833106995

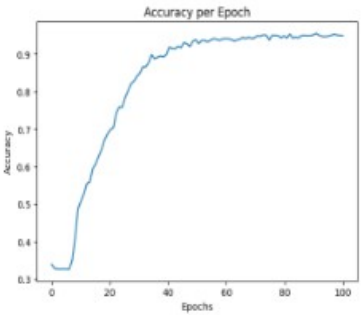

Attribute: make_up Epochs: 100 Batch-size: 100 LR: 0.003
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 68.9289071559906 seconds
Maximum Accuracy: 0.9915000200271606

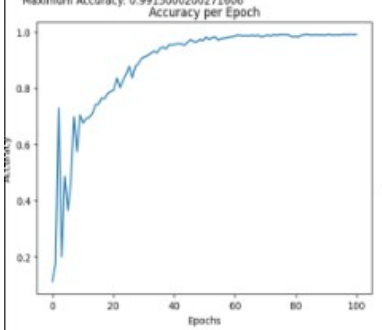

Attribute: illumination Epochs: 100 Batch-size: 100 LR: 0.005
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 103.25563287734985 seconds
Maximum Accuracy: 0.9894999861717224

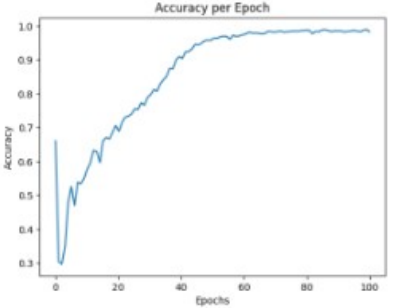

Attribute: expression Epochs: 100 Batch-size: 100 LR: 0.0055
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 64.67663717269897 seconds
Maximum Accuracy: 0.9815000295639038

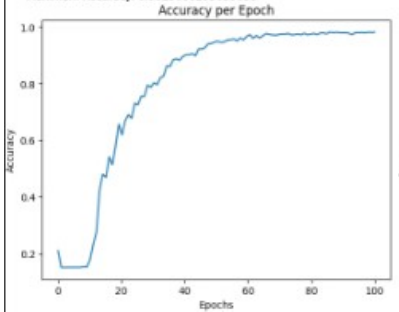

Attribute: pose Epochs: 100 Batch-size: 100 LR: 0.005
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (CPU): 142.50253820419312 seconds
Maximum Accuracy: 0.984499990940094

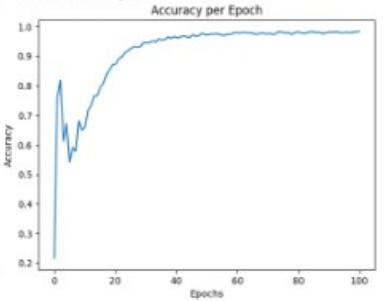

What Is a Neural Network?

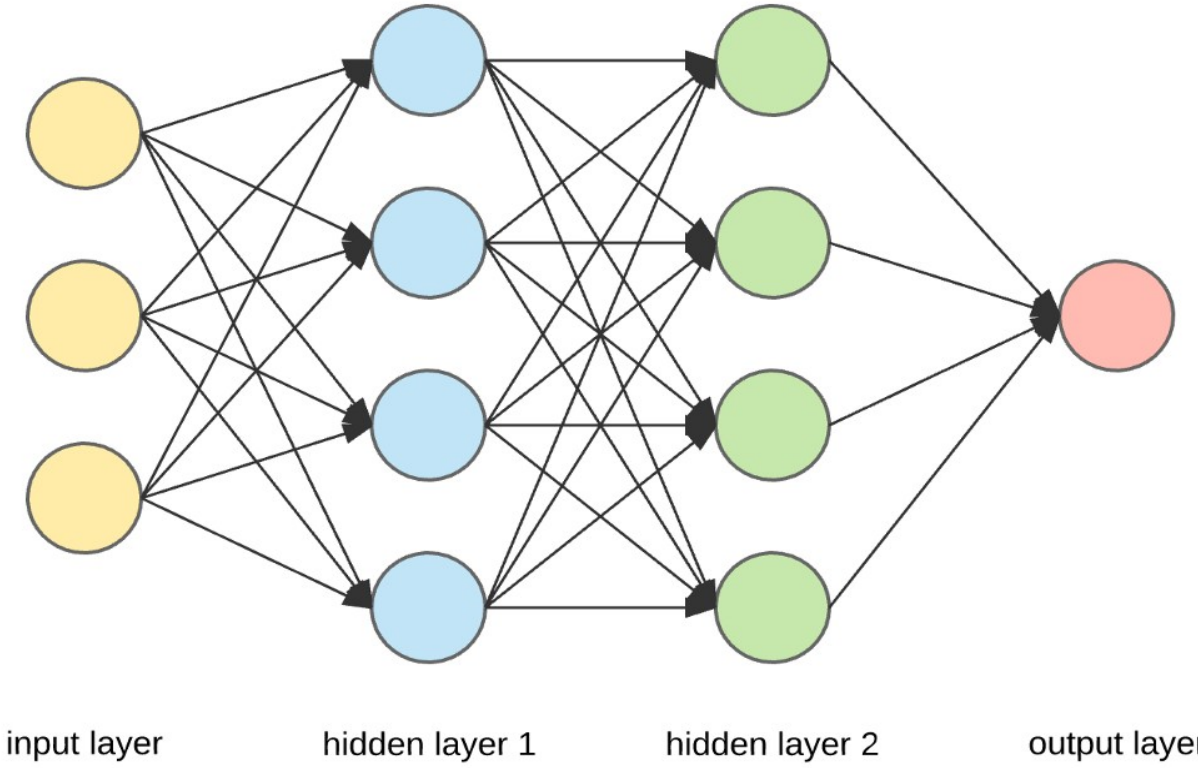

P., Petru. “What Is a Neural Network & How Does It Work? Ai Guide.” Roboflow Blog, Roboflow Blog, 14 Jan. 2025, blog.roboflow.com/what-is-a-neural-network/.

What Is a Neural Network?

Neurons

Activation

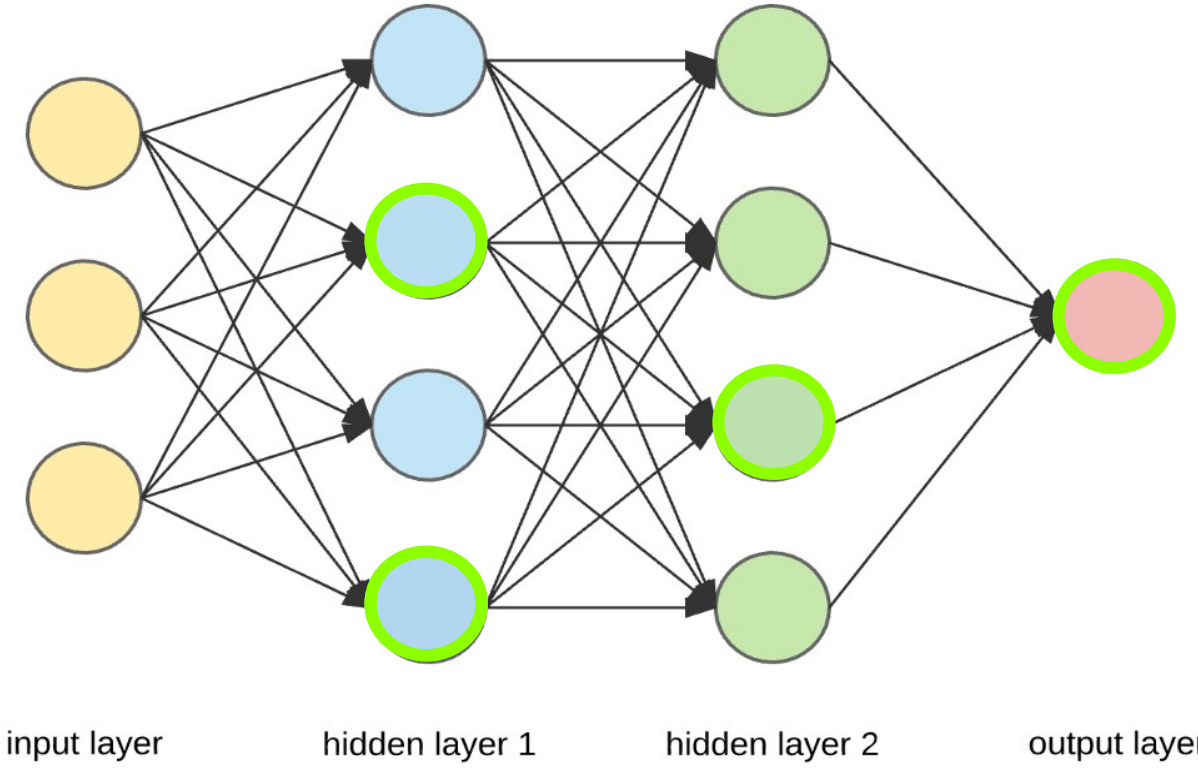

What Is a Neural Network?

Neurons

Activation

Normalization

Weights

Bias

Loss

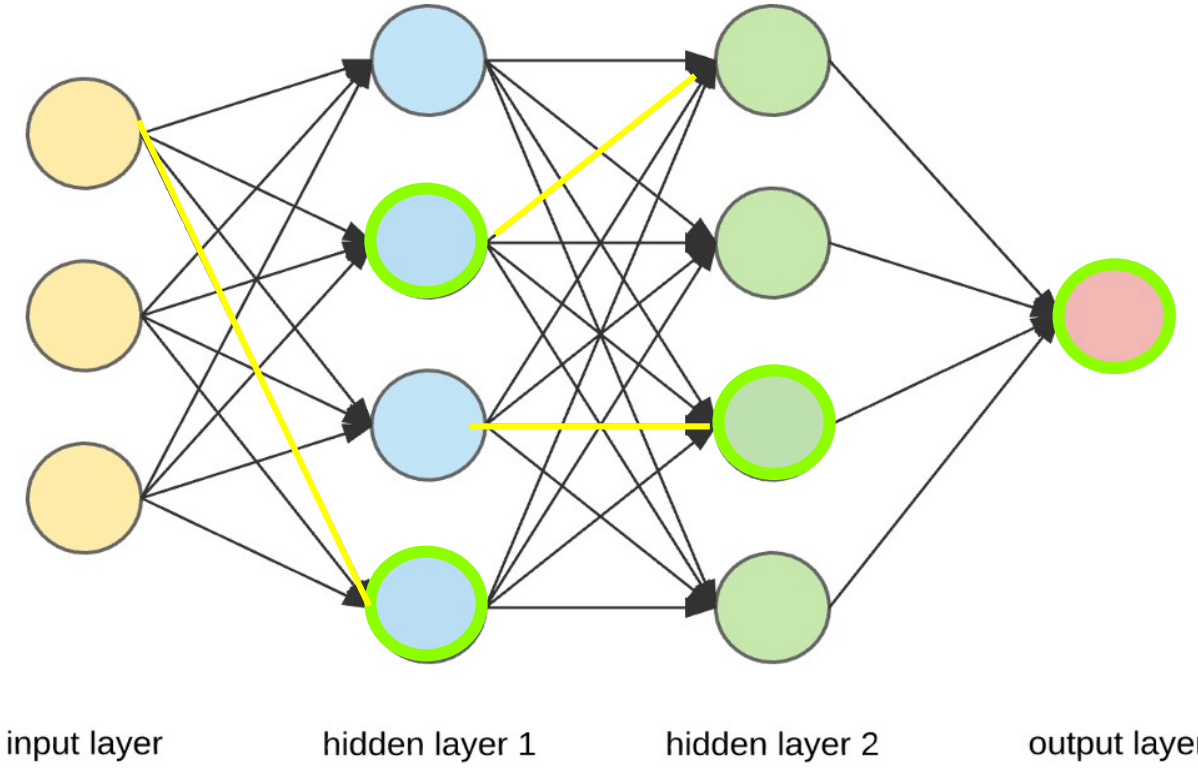

What Is a Neural Network?

Neurons

Activation

Normalization

Weights

Bias

Loss

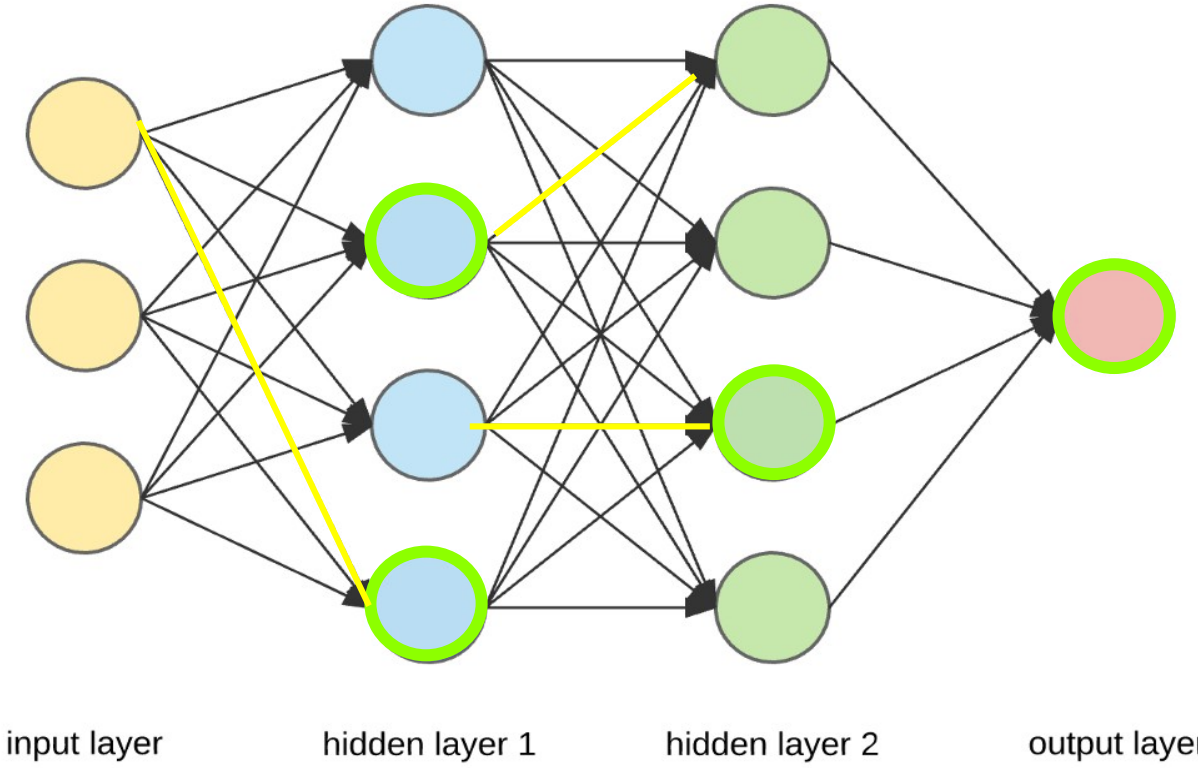

Back-propagation

Hyper-parameters:

Epochs

Learning Rate

Batch-size

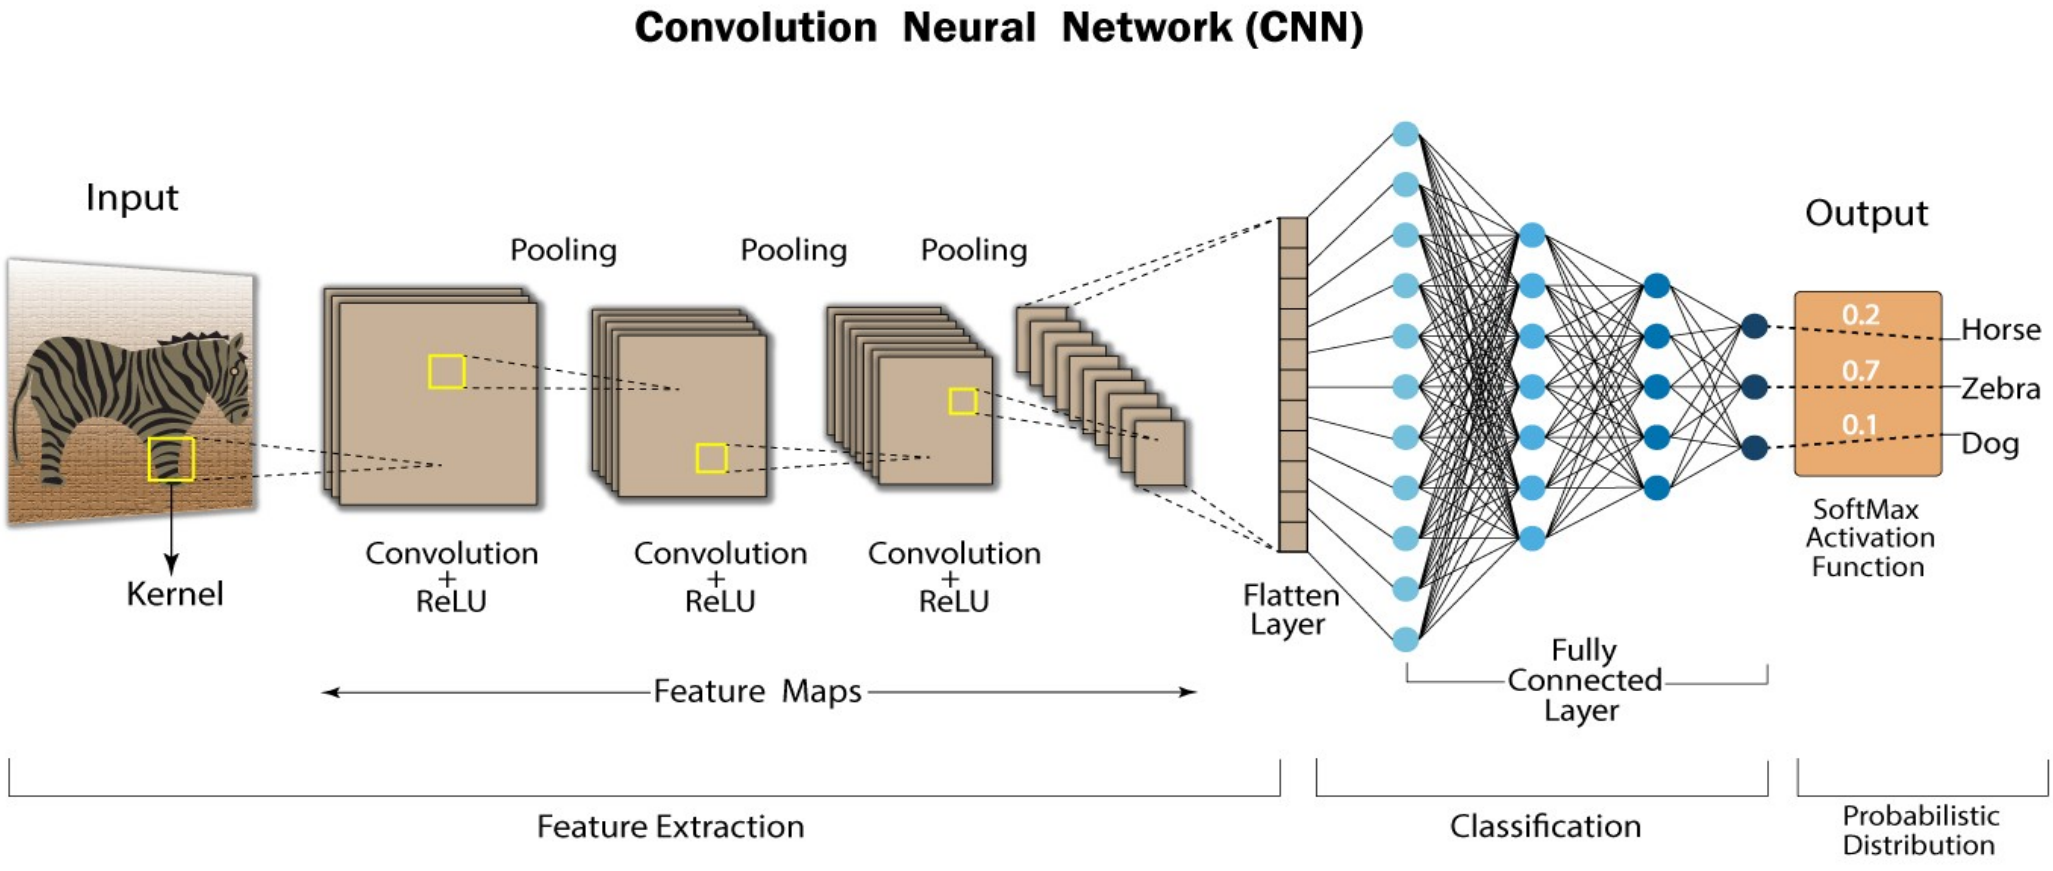

Swapna, & Sulthandudekula. (2025, February 8). Convolutional Neural Network: Deep learning. Developers Breach. https://developersbreach.com/convolution-neural-network-deep-learning/

Wider Facial Landmarks in the Wild (WFLW)
Dataset

< image1, label1>

546.262024 150.434998 543.908534 159.406036 541.577414 168.382907 539.291467 177.371376 537.070732 186.376175 534.960633 195.407389 533.198485 204.511642 532.172840 213.725502 532.149827 222
.995071 533.226431 232.200987 535.485642 241.188925 538.964011 249.778195 543.654313 257.769099 549.563517 264.902602 556.685919 270.819856 564.783499 275.317585 573.490519 278.490278 581.75
6941 280.418798 590.206692 281.109899 598.398509 279.110503 605.346333 274.284196 611.486349 268.416978 617.666489 262.591697 623.728988 256.645366 629.255988 250.203809 633.965358 243.14159
5 638.147283 235.750068 642.503118 228.461554 647.384374 221.512867 652.251029 214.554883 656.350555 207.122960 659.636538 199.294146 662.516785 191.304855 570.336914 160.796997 581.119019 1
62.115997 587.645996 168.787994 595.387939 174.869003 603.237976 179.531998 601.541992 183.760010 593.856934 179.031998 585.747986 173.287003 579.616943 166.197006 620.997986 185.745010 630.
416016 185.828003 639.794983 184.973022 648.695984 185.502014 654.565918 192.539001 646.476013 191.602020 637.629028 191.322998 628.713989 191.171005 619.406006 190.350006 606.987000 190.589
005 604.254734 202.263723 600.676132 213.660469 599.377778 224.888889 581.679016 225.307999 587.825340 229.514919 594.724702 232.292217 600.904041 233.414769 606.981934 231.963028 573.187012
177.865997 577.999123 180.307795 583.215565 181.698673 587.982847 182.864340 592.685608 184.268433 587.390062 184.756759 582.132314 184.073885 577.225492 181.613825 621.226013 195.069000 62
6.089932 195.412697 630.560012 197.345979 634.265538 199.380696 638.342346 200.496765 633.776790 200.998183 629.230190 200.430111 624.786537 198.421484 563.055054 238.871506 575.124021 239.4
17440 586.817242 242.887356 589.618218 243.393805 593.035394 244.232772 600.112982 249.218046 606.957397 254.631119 600.487464 251.939758 593.875542 249.616917 587.516509 246.675630 579.0984
11 244.212763 570.316520 243.653974 565.372009 239.214005 577.136979 241.931181 588.694857 245.436757 597.729695 248.649863 605.789307 253.858948 597.361176 248.465640 587.901542 245.213873
576.740928 241.817610 582.511111 182.722222 629.277778 200.100000 537 130 675 283 0 0 0 0 0 0 37--Soccer/37_Soccer_Soccer_37_383.jpg

Wider Facial Landmarks in the Wild (WFLW) Dataset

inates	file_name	blur	occlusion	make_up	illumination	expression	pose	y_max
95996', 5', '27...	37--Soccer/37_Soccer_soccer_ball_37_45.jpg	0	0	0	0	0	0	385
05002', 0', '14...	31--Waiter_Waitress/31_Waiter_Waitress_Waiter_...	0	0	0	0	0	0	244
93994', 0', '23...	4--Dancing/4_Dancing_Dancing_4_355.jpg	0	0	0	0	0	0	296
58640', 3', '16...	49--Greeting/49_Greeting_peoplegreeting_49_611...	0	0	0	0	0	0	360
12003', 3', '25...	29--Students_Schoolkids/29_Students_Schoolkids...	0	0	0	0	0	0	388

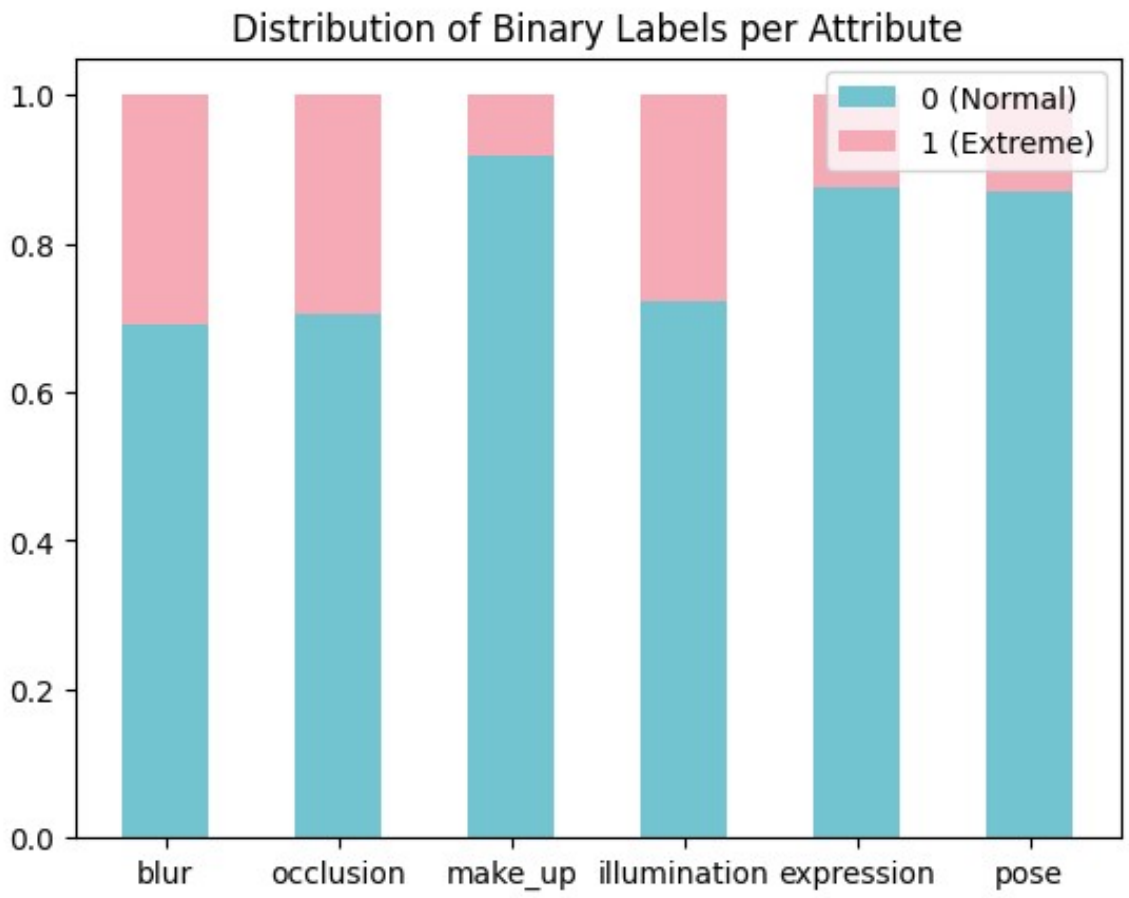
Distribution of Binary Labels per Attribute
0 (Normal)
1 (Extreme)
1.0
0.8
0.6
0.4
0.2
0.0
blur
occlusion
make_up
illumination
expression
pose

```python
model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255, input_shape=(96, 96, 3)),

    tf.keras.layers.DepthwiseConv2D(kernel_size=(3, 3), activation='relu', padding='same'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),

    tf.keras.layers.Conv2D(32, (1, 1), activation='relu', padding='same'), #changing this to increase dimension
    tf.keras.layers.DepthwiseConv2D(kernel_size=(3, 3), activation='relu', padding='same'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),

    tf.keras.layers.DepthwiseConv2D(kernel_size=(3, 3), activation='relu', padding='same'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.5),

    tf.keras.layers.Dense(1, activation='sigmoid')  # Binary classification

])
```

Epochs: 25 Batch-size: 100 Learning-rate: 0.0028

Using 2D Depthwise Convolutions

Runtime (CPU): 200.81493401527405 seconds

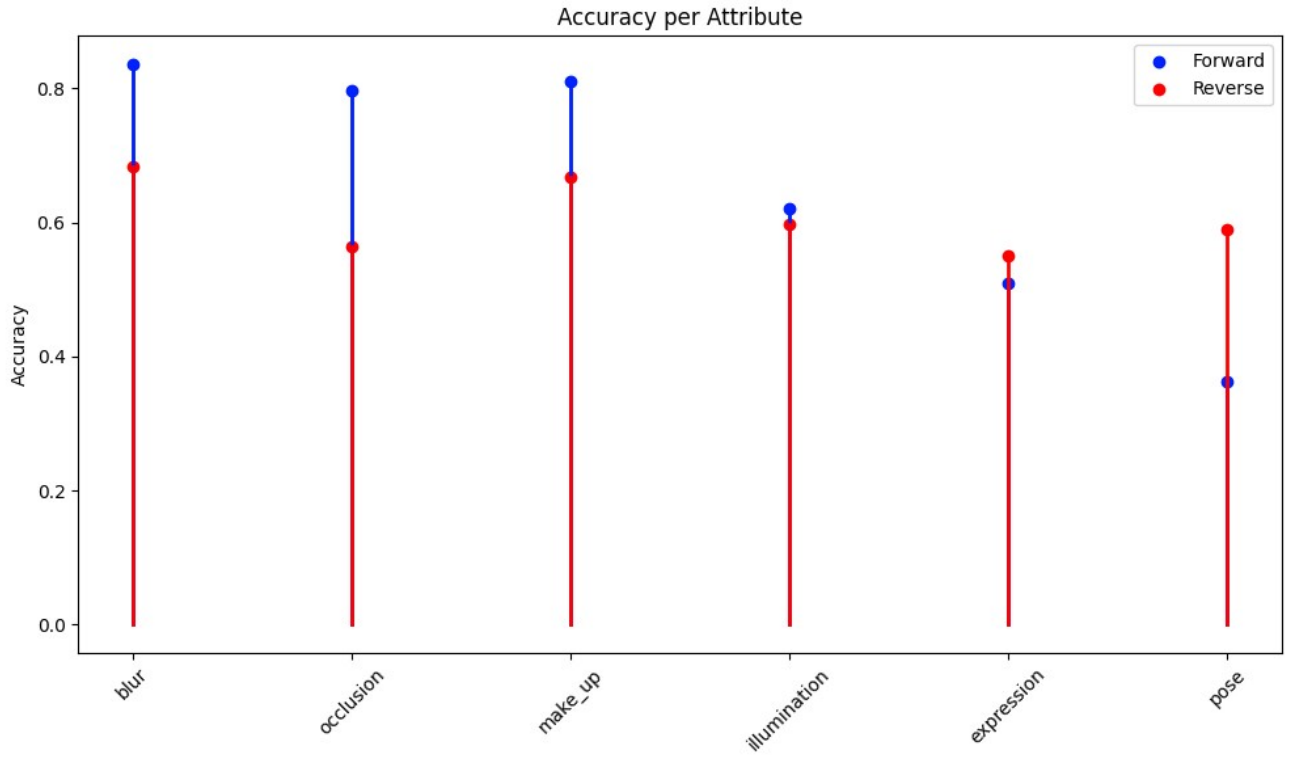

Accuracy per Epoch graph for each Attribute

Attribute: blur Epochs: 100 Batch-size: 100 LR: 0.00278
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 62.35502004623413 seconds
Maximum Accuracy: 0.9735000133514404

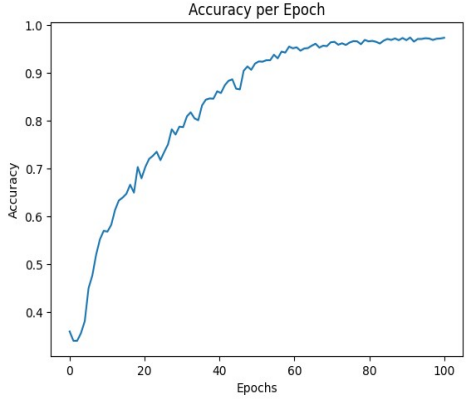

Attribute: occlusion Epochs: 100 Batch-size: 50 LR: 0.0032585
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 90.07953882217407 seconds
Maximum Accuracy: 0.9549999833106995

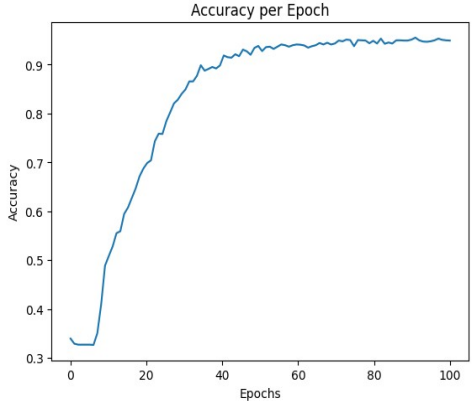

Attribute: make_up Epochs: 100 Batch-size: 100 LR: 0.003
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 68.9289071559906 seconds
Maximum Accuracy: 0.9915000200271606

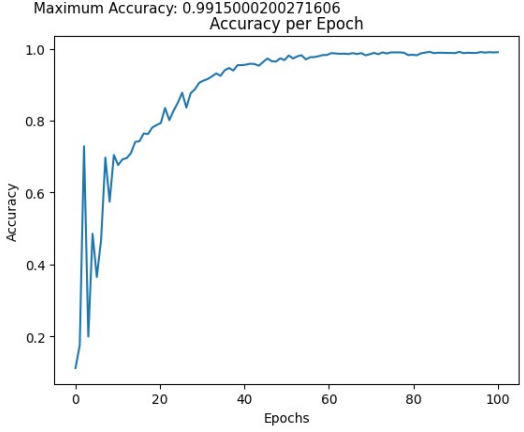

Attribute: illumination Epochs: 100 Batch-size: 100 LR: 0.005
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 103.25563287734985 seconds
Maximum Accuracy: 0.9894999861717224

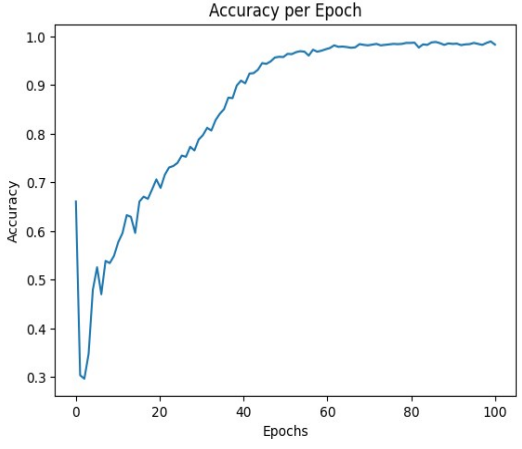

Attribute: expression Epochs: 100 Batch-size: 100 LR: 0.0055
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 64.67663717269897 seconds
Maximum Accuracy: 0.9815000295639038

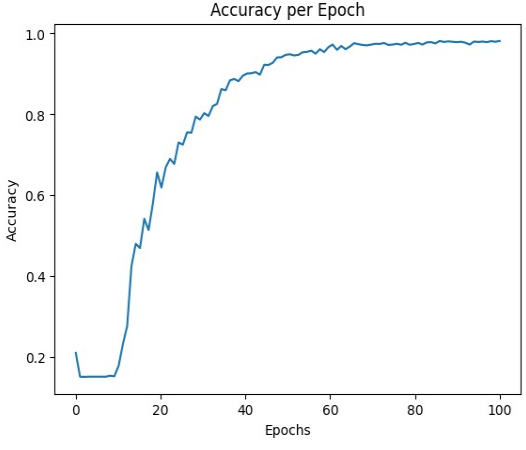

Attribute: pose Epochs: 100 Batch-size: 100 LR: 0.005
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 142.50253820419312 seconds
Maximum Accuracy: 0.984499990940094

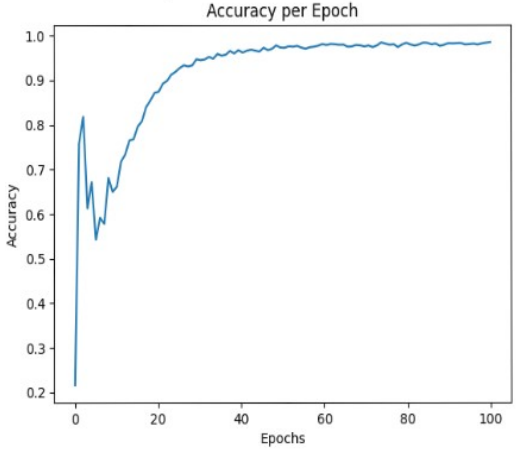

Training Results

	Blur	Occlusion	Make Up	Expression	Illumination	Pose
Accuracy	97.3%	95.5%	99.1%	98.2%	98.9%	98.4%
Learning Rate	0.00278	0.003285	0.003	0.0055	0.005	0.005
Batch-size	100	50	100	100	100	100

Testing Results

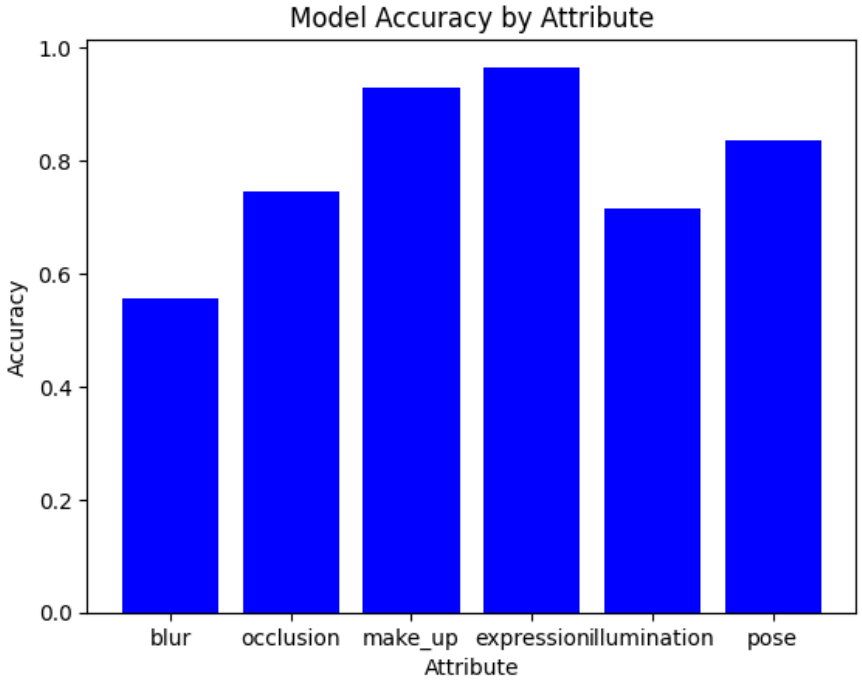

Testing Results

Extreme Illumination:

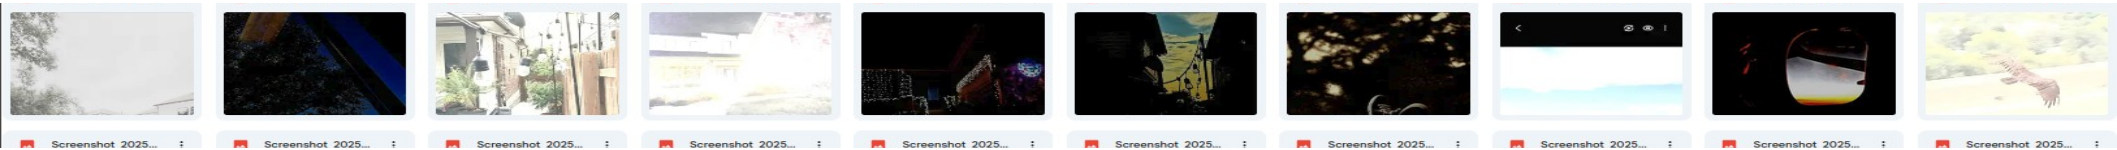

“Normal” Illumination:

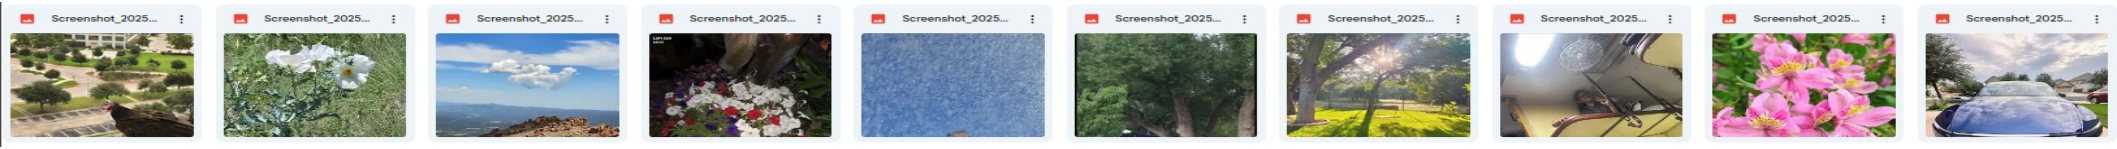

```
There are 20 testing annotations.
Loaded all 20 testing images.
1/1 ━━━━━━━━━━━━━━━━━━━━ 5s 5s/step - accuracy: 0.8500 - loss: 0.5354 - prec
illumination accuracy on my images: 0.8500000238418579
```

CPU vs GPU

Attribute: pose Epochs: 100 Batch-size: 100 LR: 0.0055
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (CPU): 25403.895089626312 seconds
Maximum Accuracy: 0.981166660785675

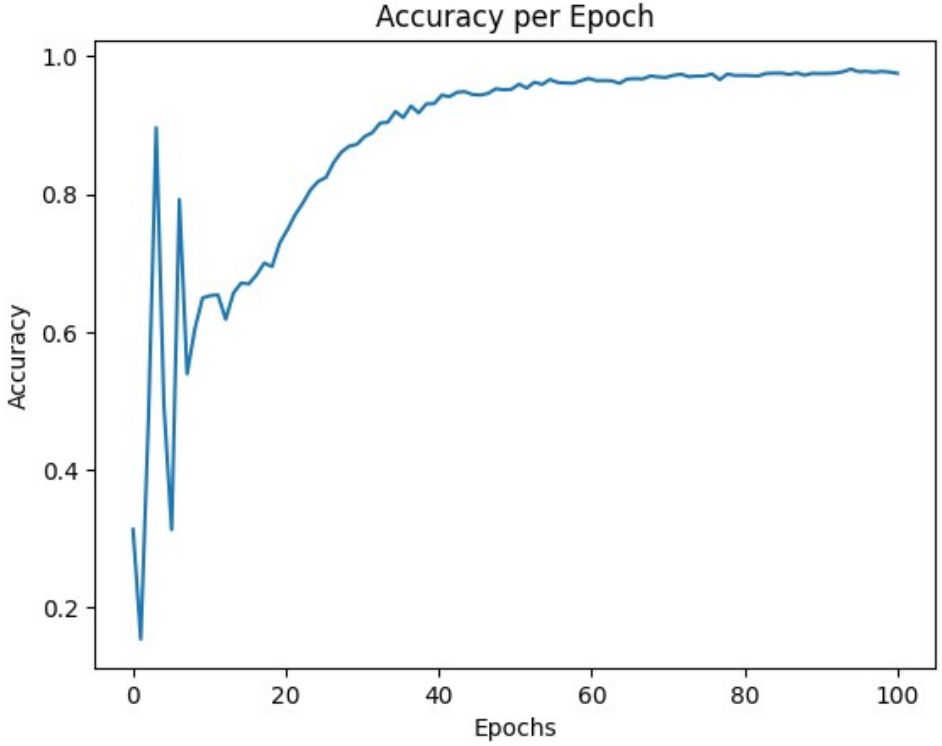

Attribute: pose Epochs: 100 Batch-size: 100 LR: 0.005
With 1 Conv2D layer between 3 DepthwiseConv2D layers
Runtime (GPU): 142.50253820419312 seconds
Maximum Accuracy: 0.984499990940094

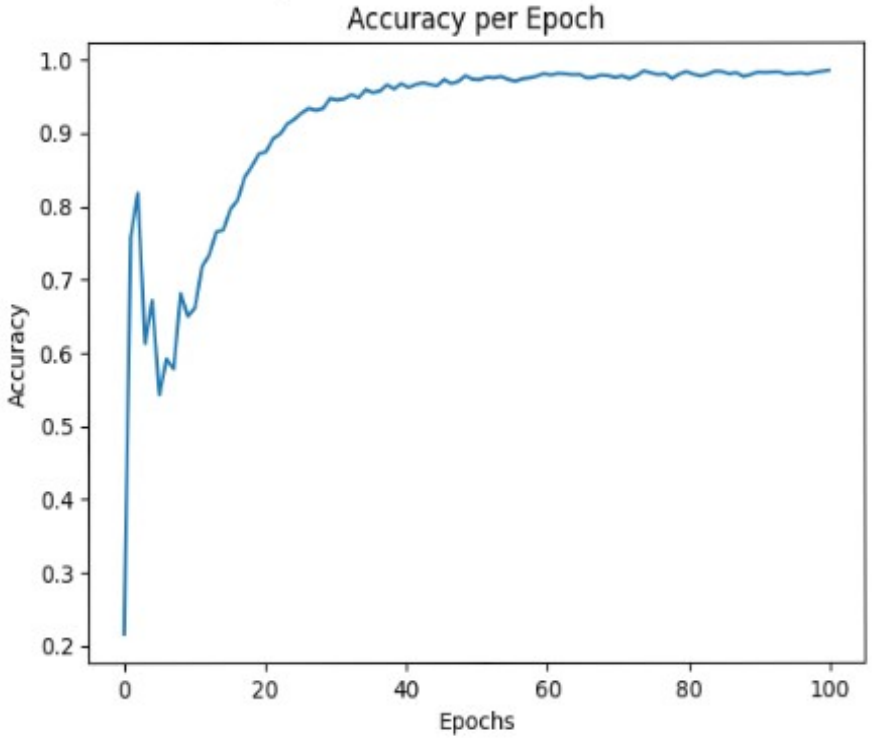

What’s Next?

Photo-Aid Application

- Manipulate structure of models
- Diverse models
- CPU time complexity reduction

Similar Applications

- Deep Fake Recognition
- Modular AI vision system

Outcomes

- Used Tensorflow, Keras, Open-CV Libraries through Google Colab

- Created Neural Networks to classify blur, occlusion, make up, illumination, expression, and pose into binary categories

- Experimented with different training conditions

- Tested using WFLW data and personal data

Thank you!

Depthwise Convolution Visual